

Package

November 27, 2018

Type Package

Title NONPARAMETRIC MLE FOR MIXTURE DISTRIBUTIONS

Version 1.1.1

Date 2018-11-24

Depends R (>= 3.0.1)

Author Nabarun Deb <nd2560@columbia.edu>
Sujayam Saha <sujayam@berkeley.edu>
Adityanand Guntuboyina <aditya@stat.berkeley.edu>
Bodhisattva Sen <bodhi@stat.columbia.edu>

Maintainer Nabarun Deb <nd2560@columbia.edu>
Sujayam Saha <sujayam@berkeley.edu>
Adityanand Guntuboyina <aditya@stat.berkeley.edu>
Bodhisattva Sen <bodhi@stat.columbia.edu>

Description Nonparametric maximum likelihood estimation in mixture models when the signal distribution is an infinite Gaussian location mixture and the signal proportion is a logistic function of the available covariate information. Other methods based on marginal likelihoods for inference in the framework of the multiple hypotheses testing problem. For a more detailed exposition, see corresponding paper: "Two-component Mixture Model in the Presence of Covariates." Nabarun Deb, Sujayam Saha, Adityanand Guntuboyina, Bodhisattva Sen. It is MANDATORY to have a moseK license in order to use this package.

License GPL-2

Encoding UTF-8

LazyData true

RoxygenNote 6.1.0.9000

URL <https://github.com/NabarunD/NPMLEmix>

Imports R.utils, dplyr, methods, latexpdf, Matrix, FDRreg, Rcpp, Rmosek, REBayes, CAMAN, progress, pbapply, Hmisc, ggplot2, pracma, mosaic

R topics documented:

NPMLEmix-package	2
makedata	4
marg1	5
marg2	7
npmleEM	8
reject_set	10

Index**12**

NPMLEmix-package	<i>Likelihood based inference in mixture models and multiple hypotheses testing</i>
------------------	---

Description

The NPMLEmix-package fits nonparametric Gaussian location mixture models for z-scores arising out of several hypotheses, while taking into account any available covariate information. It also provides three important functions: `marg1()`, `marg2()` (both based on marginal likelihoods), `npmleEM()` (based on complete data likelihood) for inference in multiple testing using local false discovery rates.

Details

Likelihoods (marginal or conditional) are computed based on the model

$$Y|X = x \sim (1 - \pi^*(x))\phi(y) + \pi^*(x) \underbrace{\int_{\theta} \phi(y - \theta) dG(\theta)}_{\phi_1(y)}, \quad X \sim m_X(\cdot)$$

where $\phi^*(\cdot)$ represents the logistic link function, $\phi(\cdot)$ is the standard Gaussian density and $G(\cdot)$ is some unknown probability measure on the real line. Usually, $\pi^*(\cdot)$ is referred to as the *signal proportion*, $\phi_1(\cdot)$ is called the *signal density* and $G(\cdot)$ is called *mixing distribution*. The i^{th} local false discovery rate is then defined as

$$lfdr_i = \frac{(1 - \pi^*(X_i))\phi_0(Y_i)}{(1 - \pi^*(X_i))\phi_0(Y_i) + \pi^*(X_i)\phi_1(Y_i)}$$

All the principal functions estimate the unknown parameters - $\pi^*(\cdot)$ and $\phi_1(\cdot)$, and consequently the $lfdr_i$'s. The optimization algorithms use quasi-Newton routines such as the BFGS (Broyden-Fletcher-Goldfarb-Shanno) algorithm and the separable convex optimization routine available in the Rmosek optimization suite. The principal functions accept a vector of z-scores (Y 's) and a covariate matrix X in their list of arguments. Read the documentations for each function to check whether or not to add a column of 1 's to X matrix.

Functions

The principal functions in the NPMLEmix-package: `marg1()`, `marg2()`, `npmleEM()`.

References

- Deb, N., Saha, S., Guntuboyina, A. and Sen, B., 2018. Two-component Mixture Model in the Presence of Covariates. arXiv preprint arXiv:1810.07897.
- Scott, J.G., Kelly, R.C., Smith, M.A., Zhou, P. and Kass, R.E., 2015. False discovery rate regression: an application to neural synchrony detection in primary visual cortex. Journal of the American Statistical Association, 110(510), pp.459-471.
- Efron, B., 2005. Local false discovery rates.
- Koenker, R. and Mizera, I., 2014. Convex optimization in R. Journal of Statistical Software, 60(5), pp.1-23.

Examples

```

require(NPMLEmix)

### Using the makedata() method ###

x=cbind(runif(1000),runif(1000))
n=1000
atoms=c(-2,0,2)
probs=c(0.48,0.04,0.48)
variances=c(1,16,1)
sx=c(-3,1.5,1.5)
### Generating the data ###
st=makedata(n,x,sx,atoms,probs,variances)
### Output the vector of local false discovery rates ###
st$localfdr

### Using the marg1() method ###
st=makedata(100,cbind(runif(100),runif(100)),c(0,1,-1),c(0,1),c(0.4,0.6),c(1,1))

### Use the default rejection level ###
defm1=marg1(st$y, cbind(1, st$xs))
### Use a new rejection level of 0.1 ###
nodefm1=marg1(st$y, cbind(1, st$xs), level = 0.1)
### Output the vector of prior probabilities ###
defm1$p
### Output the rejection set ###
nodefm1$rejset

### Using the marg2() method ###
st=makedata(100,cbind(runif(100),runif(100)),c(0,1,-1),c(0,1),c(0.4,0.6),c(1,1))

### Use the default rejection level ###
defm2=marg2(st$y, cbind(1, st$xs))
### Use a new rejection level of 0.1 ###
nodefm2=marg2(st$y, cbind(1, st$xs), level = 0.1)
### Output the vector of prior probabilities ###
defm2$p
### Output the rejection set ###
nodefm2$rejset

### Using the npmleEM() method ###
st=makedata(100,cbind(runif(100),runif(100)),c(0,1,-1),c(0,1),c(0.4,0.6),c(1,1))

### Use the default rejection level and default initialization ###
npmle1=npmleEM(st$y, cbind(1, st$xs))
### Use a new rejection level of 0.1 and marg2() initialization ###
npmle2=npmleEM(st$y, cbind(1, st$xs), level = 0.1, initp = 2)
#' ### Use a new rejection level of 0.1 and FDRreg() initialization ###
npmle3=npmleEM(st$y, cbind(1, st$xs), level = 0.1, initp = 3)
### Use the best initialization among marg1(), marg2() and FDRreg()###
npmle4=npmleEM(st$y, cbind(1, st$xs), level = 0.2, initp = 4)
### Output the vector of prior probabilities ###

```

```

npmle1$p
### Output the rejection set ###
npmle2$rejset
### Output the vector of local false discovery rates ###
npmle3$localfdr
### Output the vector of estimated conditional densities ###
npmle4$den

```

makedata

Simulates data from the aforementioned model

Description

This function can be used to simulate observations from the aforementioned model, if $G(\cdot)$ is chosen as a finite Gaussian mixture. It returns the true local false discovery rates which determine the optimal multiple testing procedure.

Usage

```
makedata(n, x, sx, atoms, probs, variances)
```

Arguments

n	Number of z-scores to be generated.
x	$n \times p$ data matrix. Do not add an additional column of 1's.
sx	The vector of coefficients for the logistic function. The first entry will be considered as the intercept term by default. Requires compatibility with x. See Details .
atoms	The vector of means for each component of the mixing distribution.
probs	The probability vector for the mixing distribution.
variances	The vector of variances for each component of the mixing distribution. Requires compatibility with atoms and probs. See Details .

Details

Given $X = x$, a $\text{Bernoulli}(\pi^*(x))$ sample is drawn. If the outcome is 1 (0), a z-score is drawn from $\phi_1(\cdot)$ ($\phi(\cdot)$). All the observations corresponding to a Bernoulli outcome 1 (0) is termed as *non-null observations* (*null observations*).

The length of sx should be 1 more than the number of columns of the data matrix x (if df = , or the transformed data matrix after a basis spline expansion).

The vectors atoms, probs and variances must have the same length.

Value

The output is a list with the following entries:

y	The vector of simulated z-scores.
x	The input data matrix.
xs	The data matrix after the columnwise basis spline expansion. Same as x if df=0.
pix	The vector of signal proportions.

f0y	The vector of standard Gaussian densities evaluated at simulated z-scores.
f1y	The vector of signal densities evaluated at simulated z-scores.
den	The vector of conditional densities evaluated at simulated z-scores.
localfdr	The vector of local false discovery rates evaluated at simulated z-scores. Note that the local FDR can be interpreted as one minus the posterior probability that a given observation is non-null.
ll	The average conditional log-likelihood.
nnind	The indices corresponding to non-null observations.

References

Basu, P., Cai, T.T., Das, K. and Sun, W., 2018. Weighted false discovery rate control in large-scale multiple testing. *Journal of the American Statistical Association*, 113(523), pp.1172-1183.

Scott, J.G., Kelly, R.C., Smith, M.A., Zhou, P. and Kass, R.E., 2015. False discovery rate regression: an application to neural synchrony detection in primary visual cortex. *Journal of the American Statistical Association*, 110(510), pp.459-471.

Examples

```
require(NPMLEmix)

x=cbind(runif(1000),runif(1000))
n=1000
atoms=c(-2,0,2)
probs=c(0.48,0.04,0.48)
variances=c(1,16,1)
sx=c(-3,1.5,1.5)
### Generating the data ###
st=makedata(n,x,sx,atoms,probs,variances)
### Output the vector of local false discovery rates ###
st$localfdr
```

marg1	<i>Implements a profile likelihood based algorithm for estimating signal proportion and density</i>
-------	---

Description

This function estimates the signal proportion and the signal density by starting from the marginal distribution of Y and followed by a profile likelihood based approach. It returns the vector of estimated local false discovery rates and the corresponding rejection set at a pre-specified level for the false discovery rate.

Usage

```
marg1(y, x, blambda = 1e-06/length(y), level = 0.05)
```

Arguments

<code>y</code>	The observed vector of z-scores.
<code>x</code>	The $n \times p$ data matrix, where n must be equal to the length of <code>y</code> . If you are interested in the intercept, you must add a column of 1's to <code>x</code> .
<code>blambda</code>	The tolerance threshold while implementing a quasi-Newton approach for estimating the signal proportion. Default is set to $1e - 6 / \text{length}(y)$. We recommend not changing it unless absolutely sure.
<code>level</code>	The level at which the false discovery rate is to be controlled. Should be a scalar in $[0, 1]$. Default set to 0.05.

Details

Note that the marginal distribution of Y based on the aforementioned model is same as that in a standard two-groups model (Efron 2008, see **References**). Fixing $\bar{\pi} = \mathbf{E}[\pi(X)]$, the signal density $\phi_1(\cdot)$ is estimated using the Rmosek optimization suite. The primary idea is to approximate the mixing distribution $G(\cdot)$ using $\max\{100, \sqrt{n}\}$ many components, each having a suitable Gaussian distribution. The signal proportion is then estimated using the BFGS algorithm. Finally, the algorithm chooses the best value of $\bar{\pi}$ based on a profile likelihood approach.

Value

This function returns a list consisting of the following:

<code>p</code>	The estimated prior probabilities, i.e., $\hat{\pi}(\cdot)$ evaluated at the data points.
<code>b</code>	The estimates for the coefficient vector in the logistic function.
<code>f1y</code>	The vector of estimated signal density evaluated at the data points.
<code>kwo</code>	This is a list with four items - i. <i>atoms</i> : The vector of means for the Gaussian distributions used to approximate $G(\cdot)$, ii. <i>probs</i> : The vector of probabilities for each Gaussian component used to approximate $G(\cdot)$, iii. <i>f1y</i> : Same as <code>f1y</code> above, iv. <i>ll</i> : The average of the logarithmic values of <code>f1y</code> .
<code>localfdr</code>	The vector of estimated local false discovery rates evaluated at the data points.
<code>den</code>	The vector of estimated conditional densities evaluated at the data points.
<code>ll</code>	The log-likelihood evaluated at the estimated optima.
<code>rejset</code>	The vector of 1's and 0's where 1 indicates that the corresponding hypothesis is to be rejected.
<code>pi0</code>	The average of the entries of the vector <code>p</code> .
<code>ll_list</code>	The vector of profile log-likelihoods corresponding to a pre-determined set of grid points for $\bar{\pi}$. The highest element of this vector is the output in <code>ll</code> .

References

- Deb, N., Saha, S., Guntuboyina, A. and Sen, B., 2018. Two-component Mixture Model in the Presence of Covariates. arXiv preprint arXiv:1810.07897.
- Koenker, R. and Mizera, I., 2014. Convex optimization, shape constraints, compound decisions, and empirical Bayes rules. Journal of the American Statistical Association, 109(506), pp.674-685.
- Efron, B., 2008. Microarrays, empirical Bayes and the two-groups model. Statistical science, pp.1-22.

Examples

```
require(NPMLEmix)
### Use example data ###
st=makedata(100,cbind(runif(100),runif(100)),c(0,1,-1),c(0,1),c(0.4,0.6),c(1,1))
### Use the default rejection level ###
defm1=marg1(st$y, cbind(1, st$xs))
### Use a new rejection level of 0.1 ###
nodefm1=marg1(st$y, cbind(1, st$xs), level = 0.1)
### Output the vector of prior probabilities ###
defm1$p
### Output the rejection set ###
nodefm1$rejset
```

marg2	<i>Implements a non-linear least squares based algorithm for estimating signal proportion and density</i>
-------	---

Description

This function estimates the signal proportion and the signal density by starting from the conditional mean $Y|X = x$ and followed by a non-linear least squares regression based approach. It returns the vector of estimated local false discovery rates and the corresponding rejection set at a pre-specified level for the false discovery rate.

Usage

```
marg2(y, x, nlslambda = 1e-06/length(y), level = 0.05)
```

Arguments

y	The observed vector of z-scores.
x	The $n \times p$ data matrix, where n must be equal to the length of y. If you are interested in the intercept, you must add a column of 1's to x.
nlslambda	The tolerance threshold while implementing a quasi-Newton approach for the non-linear least squares problem. Default is set to $1e - 6/\text{length}(y)$. We recommend not changing it unless absolutely sure.
level	The level at which the false discovery rate is to be controlled. Should be a scalar in $[0, 1]$. Default set to 0.05.

Details

Note that the conditional mean of $Y|X$ based on the aforementioned model is a non-linear function of the parameters, i.e., the logistic coefficients and the mean of the marginal distribution of Y , $\mu^* = \mathbf{E}[Y]$. This is a non-convex optimization problem in the parameters and is solved by varying μ^* over a pre-determined grid, and optimizing over the logistic coefficients. This is the estimate of $\pi^*(\cdot)$ from the marg2() method. The estimate of $\phi_1(\cdot)$ is obtained as in the marg1() method by using the Rmosek optimization suite, and the same discrete approximation to the mixing distribution $G(\cdot)$.

Value

This function returns a list consisting of the following:

<code>p</code>	The estimated prior probabilities, i.e., $\hat{\pi}(\cdot)$ evaluated at the data points.
<code>b</code>	The estimates for the coefficient vector in the logistic function.
<code>f1y</code>	The vector of estimated signal densities evaluated at the data points.
<code>kwo</code>	This is a list with four items - i. <i>atoms</i> : The vector of means for the Gaussian distributions used to approximate $G(\cdot)$, ii. <i>probs</i> : The vector of probabilities for each Gaussian component used to approximate $G(\cdot)$, iii. <i>f1y</i> : Same as <code>f1y</code> above, iv. <i>ll</i> : The average of the logarithmic values of <code>f1y</code> .
<code>localfdr</code>	The vector of estimated local false discovery rates evaluated at the data points.
<code>den</code>	The vector of estimated conditional densities evaluated at the data points.
<code>ll</code>	The log-likelihood evaluated at the estimated optima.
<code>rejset</code>	The vector of 1's and 0's where 1 indicates that the corresponding hypothesis is to be rejected.
<code>pi0</code>	The average of the entries of the vector <code>p</code> .
<code>ll_list</code>	The vector of profile log-likelihoods corresponding to a pre-determined set of grid points for μ^* . The highest element of this vector is the output in <code>ll</code> .

References

Deb, N., Saha, S., Guntuboyina, A. and Sen, B., 2018. Two-component Mixture Model in the Presence of Covariates. arXiv preprint arXiv:1810.07897.

Koenker, R. and Mizera, I., 2014. Convex optimization, shape constraints, compound decisions, and empirical Bayes rules. Journal of the American Statistical Association, 109(506), pp.674-685.

Examples

```
require(NPMLEmix)
### Use example data ###
st=makedata(100,cbind(runif(100),runif(100)),c(0,1,-1),c(0,1),c(0.4,0.6),c(1,1))
### Use the default rejection level ###
defm2=marg2(st$y, cbind(1, st$xs))
### Use a new rejection level of 0.1 ###
nodefm2=marg2(st$y, cbind(1, st$xs), level = 0.1)
### Output the vector of prior probabilities ###
defm2$p
### Output the rejection set ###
nodefm2$rejset
```

npmlEM

Implements the full likelihood approach based on the EM algorithm for estimating signal proportion and density

Description

This function estimates the signal proportion and the signal density by starting from the full likelihood of the sample and followed by an EM algorithm based approach. It returns the vector of estimated local false discovery rates and the corresponding rejection set at a pre-specified level for the false discovery rate.

Usage

```
nupleEM(y, x, level = 0.05, initp = 1)
```

Arguments

y	The observed vector of z-scores.
x	The $n \times p$ data matrix, where n must be equal to the length of y. If you are interested in the intercept, you must add a column of 1's to x.
level	The level at which the false discovery rate is to be controlled. Should be a scalar in $[0, 1]$. Default set to 0.05.
initp	The initialization method for the EM algorithm. It should be either 1, 2, 3 or 4. 1 indicates a marg1() initialization, 2 indicates a marg2() initialization, 3 indicates a FDRreg() initialization (see Details and References) and 4 chooses that initialization among marg1(), marg2() and FDRreg() which yields the highest sample likelihood. Default is set to 1.

Details

The key observation in the full likelihood approach is that the M-step of the EM algorithm results in two decoupled optimization problems, one involving $\pi^*(\cdot)$ and the other involving $\phi_1(\cdot)$. These two individual problems are then solved using the BFGS algorithm and the Rmosek optimization suite, as has been discussed previously in the **Details** sections of the methods marg1() and marg2(). The FDRreg() method was introduced in Scott et al (2015). We recommend using the version of the FDRreg() package available in https://github.com/jgscott/FDRreg/tree/master/R_pkg.

Value

This function returns a list consisting of the following:

atoms	The vector of means for the Gaussian distributions used to approximate $G(\cdot)$.
probs	The vector of probabilities for each Gaussian component used to approximate $G(\cdot)$.
f1y	The vector of estimated signal densities evaluated at the data points.
f0y	The vector of null densities evaluated at the data points.
b	The estimates for the coefficient vector in the logistic function.
p	The estimated prior probabilities, i.e., $\hat{\pi}(\cdot)$ evaluated at the data points.
ll	The log-likelihood evaluated at the estimated optima.
rejset	The vector of 1's and 0's where 1 indicates that the corresponding hypothesis is to be rejected.
den	The vector of estimated conditional densities evaluated at the data points.
localfdr	The vector of estimated local false discovery rates evaluated at the data points.

References

- Deb, N., Saha, S., Guntuboyina, A. and Sen, B., 2018. Two-component Mixture Model in the Presence of Covariates. arXiv preprint arXiv:1810.07897.
- Scott, J.G., Kelly, R.C., Smith, M.A., Zhou, P. and Kass, R.E., 2015. False discovery rate regression: an application to neural synchrony detection in primary visual cortex. Journal of the American Statistical Association, 110(510), pp.459-471.

Examples

```
require(NPMLEmix)
### Use example data ###
st=makedata(100,cbind(runif(100),runif(100)),c(0,1,-1),c(0,1),c(0.4,0.6),c(1,1))
### Use the default rejection level and default initialization ###
npmle1=npmlEM(st$y, cbind(1, st$xs))
### Use a new rejection level of 0.1 and marg2() initialization ###
npmle2=npmlEM(st$y, cbind(1, st$xs), level = 0.1, initp = 2)
#' ### Use a new rejection level of 0.1 and FDRreg() initialization ###
npmle3=npmlEM(st$y, cbind(1, st$xs), level = 0.1, initp = 3)
### Use the best initialization among other three ###
npmle4=npmlEM(st$y, cbind(1, st$xs), level = 0.2, initp = 4)
### Output the vector of prior probabilities ###
npmle1$p
### Output the rejection set ###
npmle2$rejset
### Output the vector of local false discovery rates ###
npmle3$localfdr
### Output the vector of estimated conditional densities ###
npmle4$den
```

reject_set

Finds the rejection set in a multiple testing problem

Description

This function accepts a vector of local false discovery rates from a family of hypotheses and a level parameter, to compute the rejection set.

Usage

```
reject_set(locfdr, level)
```

Arguments

locfdr	The vector of local false discovery rates (actual or estimated) corresponding to a family of hypotheses.
level	The level at which the false discovery rate is to be controlled. Should ideally be a scalar in $[0, 1]$.

Details

The problem of optimal inference in multiple hypotheses testing has been widely studied in literature. In particular, this function adopts the framework and algorithm proposed in Basu et al. See **References**.

Value

A vector of 1s and 0s with 1s indicating the hypotheses which are to be rejected.

References

- Basu, P., Cai, T.T., Das, K. and Sun, W., 2018. Weighted false discovery rate control in large-scale multiple testing. *Journal of the American Statistical Association*, 113(523), pp.1172-1183.
- Deb, N., Saha, S., Guntuboyina, A. and Sen, B., 2018. Two-component Mixture Model in the Presence of Covariates. *arXiv preprint arXiv:1810.07897*.

Examples

```
require(NPMLEmix)

x=cbind(runif(1000),runif(1000))
n=1000
atoms=c(-2,0,2)
probs=c(0.48,0.04,0.48)
variances=c(1,16,1)
sx=c(-3,1.5,1.5)
stdata=makedata(n,x,sx,atoms,probs,variances)
### Obtain the rejection set ###
reject=reject_set(stdata$lo)
```

Index

makedata, [4](#)

marg1, [5](#)

marg2, [7](#)

npmleEM, [8](#)

NPMLemix-package, [2](#)

reject_set, [10](#)